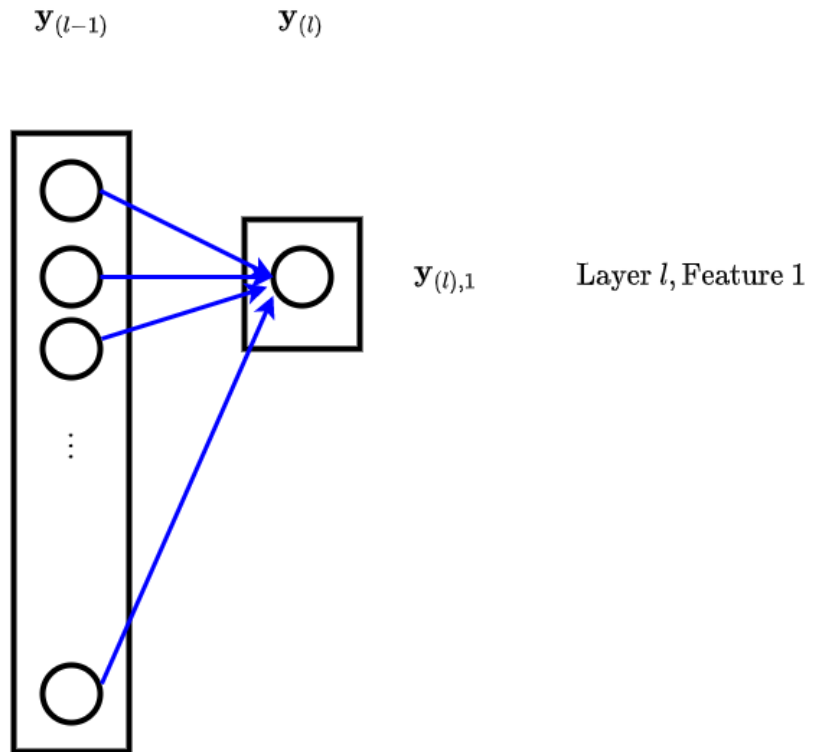


# Convolutional Neural Networks

A Fully Connected/Dense Layer with a single unit producing a single feature at layer  $l$  computes

$$\mathbf{y}_{(l),1} = a_{(l)}(\mathbf{y}_{(l-1)} \cdot \mathbf{W}_{(l),1})$$

## Fully connected, single feature



That is:

- It recognizes one new synthetic feature
- In the entirety ("fully" connected) of  $\mathbf{y}_{(l-1)}$
- Using pattern  $\mathbf{W}_{(l),1}$  (same size as  $\mathbf{y}_{(l-1)}$ )
- To reduce  $\mathbf{y}_{(l-1)}$  to a single feature.

The pattern being matched spans the entirety of the input

- Might it be useful to recognize a smaller feature that spanned only *part* of the input ?
- What if this smaller feature could occur *anywhere* in the input rather than at a fixed location ?

For example

- A "spike" in a time series
- The eye in a face

A pattern whose length was that of the entire input could recognize the smaller feature only in a *specific* place

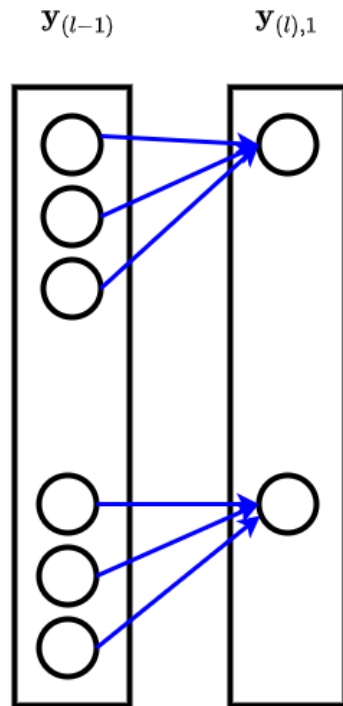
This motivates some of the key ideas behind a Convolutional Layer.

- Recognize smaller features within the whole
- Using small patterns
- That are "slid" over the entire input
- Localizing the specific part of the input containing the smaller feature

Here is the connectivity diagram of a Convolutional Layer producing a **single** feature at layer  $l$

- Using a pattern of length 3
- Eventually we will show how to produce *multiple* features
- Hence the subscript "1" in  $\mathbf{y}_{(l),1}$  to denote the first output feature
- The output  $\mathbf{y}_{(l),1}$  is called a *feature map* as it attempts to match a feature at each input location

## Convolutional layer, single feature



The important differences of a Convolutional Layer from a Fully Connected Layer:

- Produces a new *single* feature *for each location* in  $\mathbf{y}_{(l-1)}$
- $\mathbf{y}_{(l),1}$  is thus a *vector* (first feature map) of the same length as  $\mathbf{y}_{(l-1)}$
- $\mathbf{y}_{(l)}$  is a vector of  $n_{(l)}$  feature maps, one feature map per output feature
- The output feature at location  $j$  is **not** fully connected to  $\mathbf{y}_{(l-1)}$ 
  - Only a subsequence of  $\mathbf{y}_{(l-1)}$



The lack of full connectivity is significant.

In a Fully Connected network the relationship between

- Feature  $j$  and features  $(j - 1), (j + 1)$
- Is no more significant than the relationship between feature  $j$  and feature  $k \gg j$

That is: spatial locality does not matter.

To see the lack of relationship:

Let  $\text{perm}$  be a random ordering of the integers in the range  $[1 \dots n]$ .

Then

- $\mathbf{x}[\text{perm}]$  is a permutation of input  $\mathbf{x}$
- $\Theta[\text{perm}]$  is the corresponding permutation of parameters  $\Theta$ .

$$\Theta^T \cdot \mathbf{x} = \Theta[\text{perm}]^T \cdot \mathbf{x}[\text{perm}]$$

But for certain types of inputs (e.g. images) it is easy to imagine that spatial locality is important.

By using a small pattern (and restricting connectivity), we emphasize the importance of neighboring features over far way features.

Mathematically, the One Dimensional Convolutional Layer (Conv1d) we have shown computes  $\mathbf{y}_{(l)}$

$$\mathbf{y}_{(l),1} = \begin{pmatrix} a_{(l)} \left( N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, 1) \cdot \mathbf{W}_{(l),1} \right) \\ a_{(l)} \left( N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, 2) \cdot \mathbf{W}_{(l),1} \right) \\ \vdots \\ a_{(l)} \left( N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, n_{(l-1)}) \cdot \mathbf{W}_{(l),1} \right) \end{pmatrix}$$

where  $N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, j)$

- selects a subsequence of  $\mathbf{y}_{(l-1)}$  centered at  $\mathbf{y}_{(l-1),j}$

Note that

- The *same* weight matrix  $\mathbf{W}_{(l),1}$  is used for the first feature at *all* locations  $j$
- The size of  $\mathbf{W}_{(l),1}$  is the same as the size of the subsequence  $N(\mathbf{y}_{(l-1)}, \mathbf{W}_{(l),1}, j)$ 
  - Since dot product is element-wise multiplication

So  $\mathbf{W}_{(l),1}$

- Is a smaller pattern
- That is applied to *each* location  $j$  in  $\mathbf{y}_{(l-1)}$
- $\mathbf{y}_{(l),1,j}$  recognizes the match/non-match of the smaller first feature at  $\mathbf{y}_{(l-1),j}$

$\mathbf{W}_{(l),1}$  is called a convolutional *filter* or *kernel*

- We will often denote it  $\mathbf{k}_{(l),1}$
- But it is just a part of the weights  $\mathbf{W}$  of the multi-layer NN.
- We use  $f_{(l)}$  to denote the size of the smaller pattern called the *filter size*

## Note

The default activation  $a_{(l)}$  in Keras is "linear"

- That is: it returns the dot product input unchanged
- Always know what is the default activation for a layer; better yet: always specify !

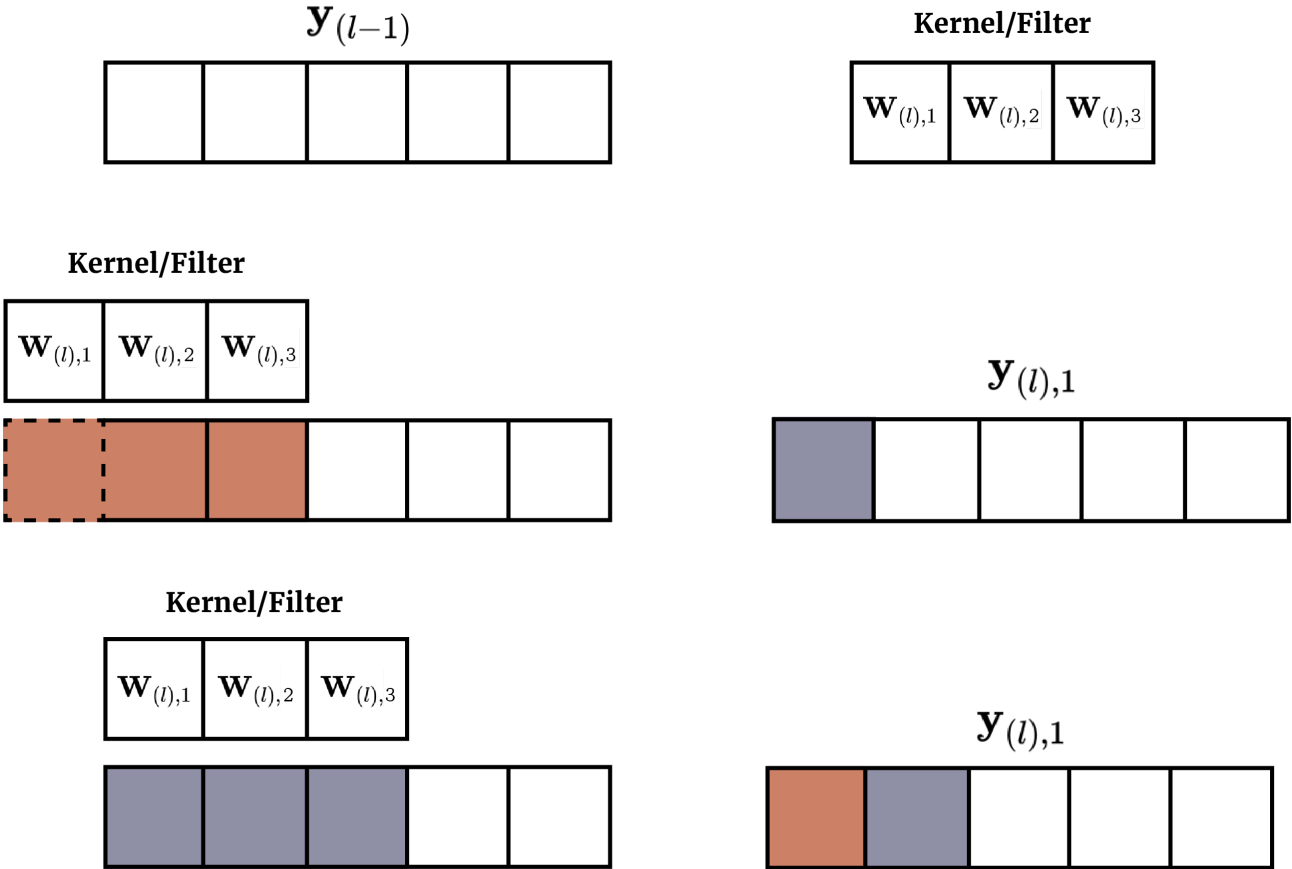


A *Convolution* is often depicted as

- A filter/kernel
- That is slid over each location in the input
- Producing a corresponding output for that location

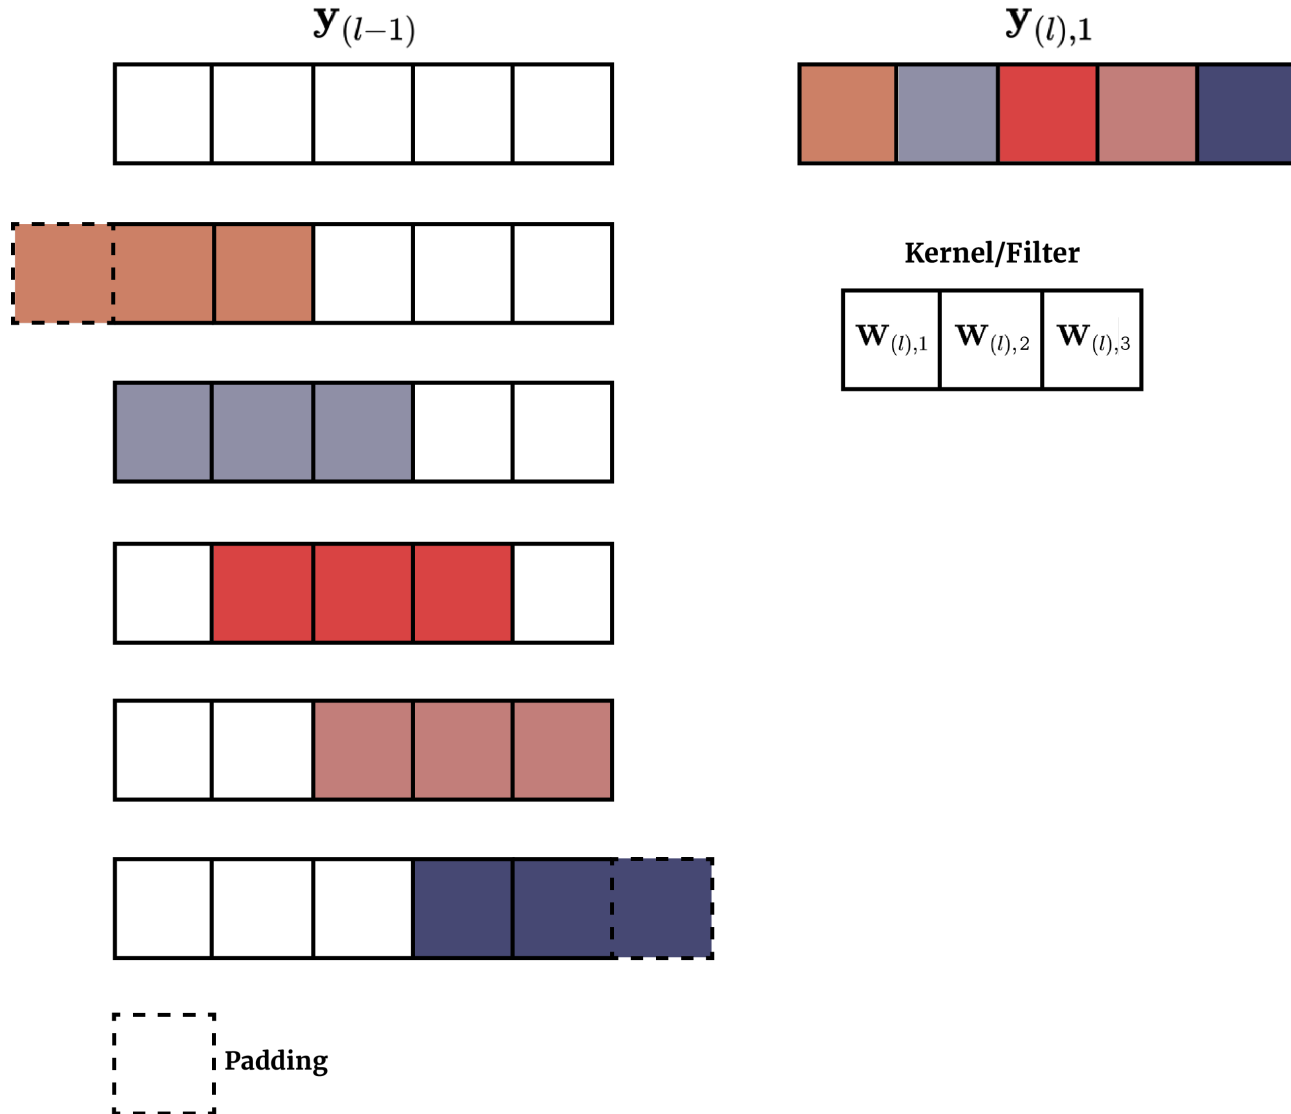
Here's a picture with a kernel of size  $f_{(l)} = 3$

Conv 1D, single feature: sliding the filter



After sliding the Kernel over the whole  $\mathbf{y}_{(l-1)}$  we get:

# Conv 1D, single feature





Element  $j$  of output  $\mathbf{y}_{(l).1}$  (i.e.,  $\mathbf{y}_{(l),1,j}$ )

- Is colored (e.g.,  $j = 1$  is colored Red)
- Is computed by applying the *same*  $\mathbf{W}_{(l),1}$  to
  - The  $f_{(l)}$  elements of  $\mathbf{y}_{(l-1)}$ , centered at  $\mathbf{y}_{(l-1),j}$
  - Which have the same color as the output

Note however that, at the "ends" of  $\mathbf{y}_{(l-1)}$  the kernel may extend beyond the input vector.

In that case  $\mathbf{y}_{(l-1)}$  may be extended with *padding* (elements with 0 value typically)

In [4]: `print("Done")`

Done