

Interpreting the coefficients in Linear Models

The dot product has been a recurring character during our Classical Machine Learning journey.

$$\Theta \cdot \mathbf{x}$$

By examining this expression more closely

- We can gain insight into what Θ means
- Understand mathematically why transformation may be necessary
- Gain an appreciation of the "pattern matching" that it embodies

Recall the places in which dot product appears

- Linear Regression

$$\hat{\mathbf{y}} = \Theta \cdot \mathbf{x}$$

- Logistic Regression

$$\hat{s} = \Theta \cdot \mathbf{x}$$

for score \hat{s} (which becomes a probability via $\hat{p} = \sigma(\hat{s})$)

- Boundary equation for linearly separable classifiers, e.g., SVM

$$0 = \Theta \cdot \mathbf{x}$$

Consider one feature $\mathbf{x}_j^{(i)}$ for example i .

- A unit increase in $\mathbf{x}_j^{(i)}$
- Holding constant the values for all other features,
- Increases $\Theta \cdot \mathbf{x}^{(i)}$ by Θ_j

Thus

$$\Theta_j = \frac{\partial}{\partial \mathbf{x}_j} \Theta \cdot \mathbf{x}$$

Θ_j may be interpreted as

- The sensitivity of $\Theta \cdot \mathbf{x}$ to changes in feature j

Numeric features

Consider numeric features $\mathbf{x}_j, \mathbf{x}_{j'}$.

Does

$$\Theta_j > \Theta_{j'}$$

mean that feature j is "more important" than feature j' ?

- No !
- It just means it has a larger impact
- Which can *also* occur if $\mathbf{x}_j, \mathbf{x}_{j'}$ are on different scales

For example consider the equality

$$\mathbf{y} = \Theta \cdot \mathbf{x}$$

- Replacing \mathbf{x}_j
- By $\mathbf{x}_{j''} = \mathbf{x}_j * 10$
- Mathematically results in $\Theta_{j''} = \Theta_j / 10$

Thus, the scale of the parameter is dependent on the scale of the feature.

Unless two features are on the same scale: we can't directly compare their corresponding parameters.

Categorical features

Consider a categorical feature with categories from

$$C = \{c_1, c_2, \dots\}$$

One Hot Encoding this feature replaces the original feature with $\|C\|$ binary features

- Is_{c_1}
- Is_{c_2}
- \vdots
- $\text{Is}_{c_{\|C\|}}$

Suppose \mathbf{x}_j corresponds to the binary feature

$$\mathbf{Is}_{c_1}$$

Then, by the formula for dot product

- Θ_j is the *increment* to $\Theta \cdot \mathbf{x}$
- Arising from $\mathbf{x}_j^{(i)} = 1$
- Compared to $\mathbf{x}_j^{(i)} = 0$

That is:

- Θ_j is how much $\Theta \cdot \mathbf{x}$ increases
- When example i has feature value c_1 rather than any of $\{c_2, \dots, \}$

We can use this interpretation

- To further emphasize the problem of
- Treating a categorical variables as a number rather than a collection of binary indicator variables

For example, let's revisit the Passenger Class $Pclass \in \{1, 2, 3\}$ from the Titanic example.

- As a collection of binary indicator variables, the increment of being in each class is $\Theta_{Is_1}, \Theta_{Is_2}, \Theta_{Is_3}$
- As a numeric variable with parameter value Θ_j
 - Being in Class 3 has three times the effect as being in Class 1

Thus

- As numeric, we imply a particular magnitude with each category
- As binary indicator, the magnitude is determined by the data

Motivating a transformation

Understanding the meaning of Θ may help us choose a transformation.

Suppose we have examples $\langle \mathbf{X}, \mathbf{y} \rangle$ where

$$\mathbf{y} = \Theta \cdot \mathbf{x}$$

does not seem to hold.

Perhaps a transformation to either/both of \mathbf{x} , \mathbf{y} **will** make the relationship linear.

For example, consider

- \mathbf{y} and \mathbf{x} are time-series of prices, with different scales
- We observe that the impact on \mathbf{y} of a *unit change* in \mathbf{x}_j
 - Is much bigger when \mathbf{x}_j is small
 - Compared to when \mathbf{x}_j is large
- So

$$\mathbf{y} \neq \Theta \cdot \mathbf{x}$$

Now suppose we re-dominate (by transforming) the timeseries \mathbf{y} and \mathbf{x} to

- Timeseries \mathbf{y}'
= daily % change in \mathbf{y}
- Timeseries \mathbf{x}'
= daily % change in \mathbf{x}

This transformation from Price to Return is common in Finance.

It may now turn out that

- A unit change in \mathbf{x}'
- Results in a change in \mathbf{y}'
- That is *independent* of the magnitude of \mathbf{x}'
- So

$$\mathbf{y}' = \Theta' \cdot \mathbf{x}'$$

That is

- A 1 percentage point change in the price of \mathbf{x}
- Causes \mathbf{y} to change by Θ' percentage points

So there **was** a relationship between \mathbf{y} and \mathbf{x} , not in Price but in Return.

The Capital Asset Pricing Model of Finance postulates such a relationship and it is common to transform prices to returns.

Transformed targets

At times we may apply transformations to target values rather than just features.

This means that Θ_j is the sensitivity of the *transformed* target.

Recall that Logistic Regression could be formulated as

- Linear Regression of the features
- Versus the *log odds*

$$\log_e \frac{\hat{p}}{1 - \hat{p}} = \Theta^T \mathbf{x}$$

So a unit change in feature \mathbf{x}_j with parameter Θ_j changes the odds $\frac{\hat{p}}{1-\hat{p}}$ in a *multiplicative* way

$$\log\left(\frac{\hat{p}}{1-\hat{p}}\right) + \Theta_j = \log\left(\frac{\hat{p}}{1-\hat{p}} * \exp \Theta_j\right)$$

-

Examples

- Log transform of target:
 - $\log \mathbf{y} = \Theta_0 + \Theta_1 * \mathbf{x}_1$
 - $\theta_1 = \frac{\partial \log \mathbf{y}}{\partial \mathbf{x}_1} = \% \text{ change in } \mathbf{y} \text{ per unit change in } \mathbf{x}_1$
- Log transform of both target and feature:
 - $\log \mathbf{y} = \Theta_0 + \Theta_1 * \log \mathbf{x}_1$
 - $\Theta_1 = \frac{\partial \log \mathbf{y}}{\partial \log \mathbf{x}_1} = \% \text{ change in } \mathbf{y} \text{ per } \% \text{ change in } \mathbf{x}_1$
- Standardize feature
 - Transform \mathbf{x} into $z_{\mathbf{x}} = \frac{\mathbf{x} - \bar{\mathbf{x}}}{\sigma_{\mathbf{x}}}$
 - $\mathbf{y} = \Theta_0 + \Theta_1 * z_{\mathbf{x}}$
 - $\Theta_1 = \frac{\partial \log \mathbf{y}}{\partial z_{\mathbf{x}}}$ change in \mathbf{y} per 1 standard deviation change in \mathbf{x}
 - since z is in units of "number of standard deviations"

Remember

- if you transform features in training, you must apply the same transformation to features in test
 - if the transformation is parameterized, the parameters are determined at **train** fit time, not test !
- if you transform the target, the prediction is in different units than the original
 - you can perform the inverse transformation to get a prediction in original units

Bucketing/Binning re-visited

Suppose x_j is a continuous numeric feature (e.g., Age).

Some questions to consider

- Is a 1 year increase in age equally relevant for all ages ?
 - If so: numeric
- Is a 1 year increase in age of the same relevance for a senior adult compared to an infant ?
 - If not: consider reducing discrete ages to discrete buckets
 - Is there a linear relationship between target and the center point of the bucket ?
 - If so: bucket feature can be numeric
 - If not: bucket feature categorical

Interpreting the MNIST classifier: template matching

The Θ produced by a linear classifier can be viewed as templates

- the strength of Θ_j tells you how strongly feature \mathbf{x}_j influences the target

So we can interpret Θ as a "template" for what a model is looking for.

Let's look at the template for

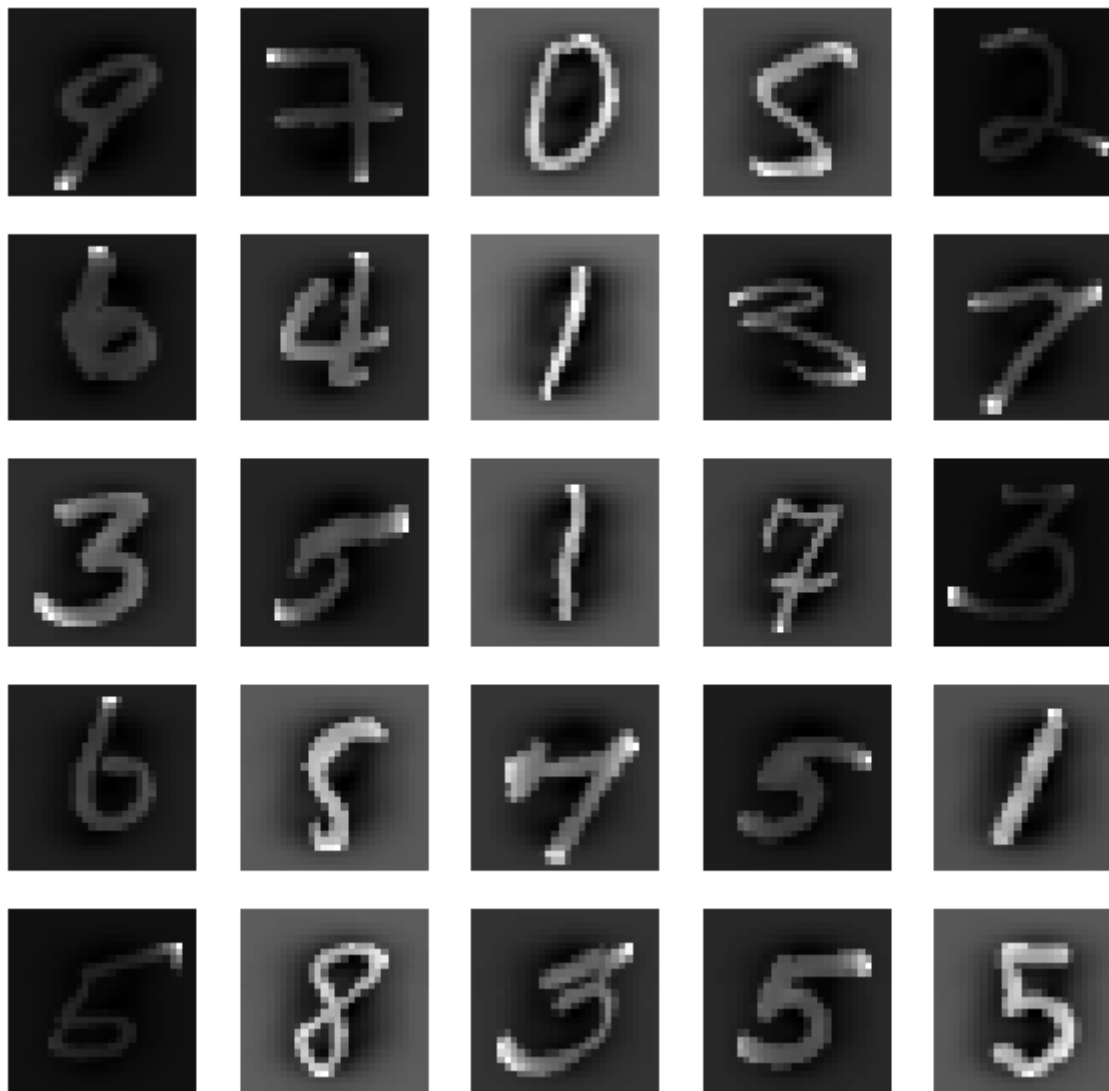
- The 10 separate, single-digit binary MNIST classifiers
- Or similarly: each row of Θ for the multinomial 10 class MNIST classifier

Here's the training data


```
In [5]: mnh.setup()  
        mnh.visualize()
```

Retrieving MNIST_784 from cache

Out[5]:






Let's fit a `LogisticRegression` model and examine the templates (coefficients Θ)

```
In [6]: _ = mnh.fit()  
mnist_fig, mnist_ax = mnh.plot_coeff()
```

Parameters for...



Recall

- There is one parameter per pixel
- The parameters are ordered in the same way as the linearization of the pixels
 - from (28×28) grid to a vector of 784 numbers.
- We can display the 784 parameters in a (28×28) image to show the intensity of parameter associated with a pixel
- White is high parameter value; Black is low (or negative)

- The template for 0 emphasizes small values (absence of bright pixels) in the center of the image
- The template for 1 emphasizes bright vertical pixels
- The template for 8 emphasizes the absence of bright pixels
 - in the two circles
 - in the pinched waist

You can now imagine how these templates might lead to misclassification

What is the classification of

- a "7" with a strong vertical line in the center (that's what the "1" template tries to match)
- a thin "0" (the "0" template is looking for a large donut)

So interpretation is a very powerful diagnostic tool for both understanding and improving your models.

In [7]: `print("Done")`

Done