Inside the LSTM: update equations

An LSTM layer, at time step t

- Takes input element $\mathbf{x}_{(t)}$
- Updates long term memory $\mathbf{c}_{(t)}$
- Updates control state $\mathbf{h}_{(t)}$
- Optionally outputs $\mathbf{y}_{(t)}$

according to the equations

$$\mathbf{c}_{(t)} \;=\; \mathbf{remember}_{(t)} \otimes \mathbf{c}_{(t-1)} + \mathbf{save}_{(t)} \otimes \mathbf{c}'_{(t)} \quad \text{Long term memory}$$

$$\mathbf{h}_{(t)} \;=\; \mathbf{focus}_{(t)} \otimes \tanh(\mathbf{c}_{(t)}) \qquad\qquad\qquad \text{Short term memory/contro}$$

$$\mathbf{y}_{(t)} \;=\; \mathbf{h}_{(t)} \qquad\qquad\qquad\qquad\qquad\qquad \text{Output}$$

where

remember is a gate that allows elements of $\mathbf{c}$ to be remembered/forgotten

focus is a mask that controls movement from long-term to short-term

save is a gate that controls selective updating of elements of $\mathbf{c}$

A lot of moving parts !

The important thing to remember is that the layer

- Has a matrix $\mathbf{W}$ of weights
- That controls the update of each part

Training, as usual, seeks to discover the optimal (i.e., loss function minimizing) values for $\mathbf{W}$

Let's try to understand the whole by examining each piece.

Memory/States

$\mathbf{c}_{(t)}$ is the long-term memory.

It is a vector of features that need to be retained throughout the computation

- As each element $\mathbf{x}_{(t)}$ of input sequence $\mathbf{x}$ is processed
- It records the "concepts" that are important to solve the task
- Might have many elements

$\mathbf{h}_{(t)}$ is the short-term or control memory.

Very much like a vanilla RNN, it's job is to guide the transition from state $\mathbf{h}_{(t)}$ to state $\mathbf{h}_{(t+1)}$.

It is a vector of features that need to be retained *for the immediate future*

- It might have many elements
- Same length as $\mathbf{c}_{(t)}$

An analogy might help.

Suppose you are driving a car in an unfamiliar city

- Short term memory is a map of the surrounding blocks
- Long term memory is a map of the city plus rules of the road

Gates/Masks

remember, **save**, **focus** are vectors that interact with long term memory $\mathbf{c}_{(t)}$

- Element-wise
- So have same length as $\mathbf{c}_{(t)}$

They will be used

- To selectively modify individual elements of $\mathbf{c}_{(t)}$
- Forget/Reset the value of an element that is no longer relevant
- Decide which individual elements to update

The [classic paper that introduced the LSTM (https://www.bioinf.jku.at/publications/older/2604.pdf)](https://www.bioinf.jku.at/publications/older/2604.pdf) gives these gates different names

- $\mathbf{remember}_{(t)} \mapsto \mathbf{f}_{(t)}$
    - $\mathbf{f}$ denotes "Forget" (although it really means "don't forget", i.e, remember !)
- $\mathbf{save}_{(t)} \mapsto \mathbf{i}_{(t)}$
    - $\mathbf{i}$ denotes "Input"
- $\mathbf{focus}_{(t)} \mapsto \mathbf{o}_{(t)}$
    - $\mathbf{o}$ denotes "Output"

Hopefully the names in our presentation add clarity.

Output

$\mathbf{y}_{(t)}$ is the value (if any) output at step t, for

- A one to many function
- Or a many to many function

As written

$$\mathbf{y}_{(t)} = \mathbf{h}_{(t)}$$

so it has the same length as a memory element $\mathbf{h}_{(t)}, c_{(t)}$.

This assumption is purely for simplicity

- You can map $\mathbf{h}_{(t)}$ through another layer
- That transforms $\mathbf{h}_{(t)}$ into the appropriate type/shape for output $\mathbf{y}_{(t)}$

In fact, our equation for the vanilla RNN included this final transformation of $\mathbf{h}$ to $\mathbf{y}$:

$$\mathbf{y}_{(t)} = \mathbf{W}_{hy}\mathbf{h}_{(t)} + \mathbf{b}_y$$

The update process

Let's examine the update equation for each of the parts.

Update long-term memory

Long-term memory is updated in a two step process

- Produce a "candidate" updated value for each element of the state
- Decide which of the candidate updated values get applied to the long term memory
    - Successful candidates become part of long term memory
    - Unsuccessful candidates are dropped

The candidate update value vector $\mathbf{c}'_{(t)}$ is a function of

- The prior short term state $\mathbf{h}_{(t-1)}$
- And the current input $\mathbf{x}_{(t)}$
- Controlled by parts of the weight matrix $\mathbf{W}$

$$\mathbf{c}'_{(t)} = \tanh(\mathbf{W}_{x,c}\mathbf{x}_{(t)} + \mathbf{W}_{h,c}\mathbf{h}_{(t-1)} + \mathbf{b}_c)$$

This is very much like the RNN state update equation

- Although the RNN equation has $\mathbf{h}$ on both sides of the equation, so is directly recursive in form

We now need to decide which elements of $\mathbf{c}_{(t)}$ to change.

The **remember** mask controls forgetting the current value $\mathbf{c}_{(t-1)}$

- When $\mathbf{remember}_{(t),j} = 0$
    - $\mathbf{c}_{(t-1),j}$, the j^{th} element of $\mathbf{c}_{(t-1)}$
    - Will be reset to 0 ("forgotten")
- When $\mathbf{remember}_{(t),j} = 1$
    - $\mathbf{c}_{(t-1),j}$, the j^{th} element of $c_{(t-1)}$
    - Will contribute to the new value $\mathbf{c}_{(t),j}$

The $\mathbf{save}$ mask controls whether the candidate value $\mathbf{c}'_{(t)}$ contributes to the new value $\mathbf{c}_{(t)}$:

- When $\mathbf{save}_{(t),j} = 1$
    - Candidate value $\mathbf{c}'_{(t),j}$ will contribute to the new value $\mathbf{c}_{(t),j}$
- When $\mathbf{save}_{(t),j} = 0$
    - Candidate value $\mathbf{c}'_{(t),j}$ will **not** contribute to the new value $\mathbf{c}_{(t),j}$

Here is the update equation for $\mathbf{c}_{(t)}$.

- It combines the remember/forget decision for each element
- With the decision on passing through the candidate value for the element

$$\mathbf{c}_{(t)} = \mathbf{remember}_{(t)} \otimes \mathbf{c}_{(t-1)} + \mathbf{save}_{(t)} \otimes \mathbf{c}'_{(t)}$$

The role of $\tanh$ in the candidate value equation

Why modify the candidate value $\mathbf{c}'_{(t)}$ by passing it through $\tanh$?

The $\tanh$ has the important property

- That its range is $[-1, +1]$

So updates to $\mathbf{c}_{(t)}$ have the flavor of either

- Incrementing existing value $\mathbf{c}_{(t-1),j}$ by 1
- Or decrementing existing value $\mathbf{c}_{(t-1),j}$ by 1

This makes $\mathbf{c}_{(t)}$ act like a *counter*.

Update short-term memory (control state)

The short-term memory update

- Selectively copies parts of the newly updated long-term memory $\mathbf{c}_{(t)}$
- To short-term memory
$$\mathbf{h}_{(t)} \;=\; \mathbf{focus}_{(t)} \otimes \tanh(\mathbf{c}_{(t)}) \quad \text{Short term memory/control}$$

The $\mathbf{focus}$ mask selects which elements of long-term memory are immediately relevant for control

The $\tanh$ activation function applied to long-term memory $\mathbf{c}_{(t)}$

- Squashes the range

The gate/mask update equations

All of the gates are updated via similar equations, taking

- The prior short term state $\mathbf{h}_{(t-1)}$
- And the current input $\mathbf{x}_{(t)}$
- Controlled by parts of the weight matrix $\mathbf{W}$

$$\textbf{remember}_{(t)} \quad = \quad f_{(t)} \quad = \quad \sigma(\mathbf{W}_{x,f}\mathbf{x}_{(t)} + \mathbf{W}_{h,f}\mathbf{h}_{(t-1)} + \mathbf{b}_f)$$

$$\textbf{save}_{(t)} \quad\quad = \quad i_{(t)} \quad = \quad \sigma(\mathbf{W}_{x,i}\mathbf{x}_{(t)} + \mathbf{W}_{h,i}\mathbf{h}_{(t-1)} + \mathbf{b}_i)$$

$$\textbf{focus}_{(t)} \quad\quad = \quad o_{(t)} \quad = \quad \sigma(\mathbf{W}_{x,o}\mathbf{x}_{(t)} + \mathbf{W}_{h,o}\mathbf{h}_{(t-1)} + \mathbf{b}_o)$$

Notice the use of the sigmoid activation for each gate/mask ?

- This restricts the range of each element to $[0, 1]$
- As needed by a gate/mask

Conclusion

That was quite a workout.

There were lots of moving parts, but hopefully you can now understand each.

To conclude, here is the full set of update equations

$$
\begin{aligned}
\mathbf{c}'_{(t)} &= \tanh(\mathbf{W}_{x,c}\mathbf{x}_{(t)} + \mathbf{W}_{h,c}\mathbf{h}_{(t-1)} + \mathbf{b}_c) && \text{Candidate updat} \\
\mathbf{c}_{(t)} &= \mathbf{remember}_{(t)} \otimes \mathbf{c}_{(t-1)} + \mathbf{save}_{(t)} \otimes \mathbf{c}'_{(t)} && \text{Long term memo} \\
\mathbf{h}_{(t)} &= \mathbf{focus}_{(t)} \otimes \tanh(\mathbf{c}_{(t)}) && \text{Short term memo} \\
\mathbf{y}_{(t)} &= \mathbf{h}_{(t)} && \text{Output}
\end{aligned}
$$

where

$$
\begin{aligned}
\mathbf{remember}_{(t)} &= \sigma(\mathbf{W}_{x,f}\mathbf{x}_{(t)} + \mathbf{W}_{h,f}\mathbf{h}_{(t-1)} + \mathbf{b}_f) \\
\mathbf{save}_{(t)} &= \sigma(\mathbf{W}_{x,i}\mathbf{x}_{(t)} + \mathbf{W}_{h,i}\mathbf{h}_{(t-1)} + \mathbf{b}_i) \\
\mathbf{focus}_{(t)} &= \sigma(\mathbf{W}_{x,o}\mathbf{x}_{(t)} + \mathbf{W}_{h,o}\mathbf{h}_{(t-1)} + \mathbf{b}_o)
\end{aligned}
$$

$$
\begin{aligned}
\mathbf{remember}_{(t)} \quad &= \quad f_{(t)} \quad = \quad \sigma(\mathbf{W}_{x,f}\mathbf{x}_{(t)} + \mathbf{W}_{h,f}\mathbf{h}_{(t-1)} + \mathbf{b}_f) \\
\mathbf{save}_{(t)} \quad &= \quad i_{(t)} \quad = \quad \sigma(\mathbf{W}_{x,i}\mathbf{x}_{(t)} + \mathbf{W}_{h,i}\mathbf{h}_{(t-1)} + \mathbf{b}_i) \\
\mathbf{focus}_{(t)} \quad &= \quad o_{(t)} \quad = \quad \sigma(\mathbf{W}_{x,o}\mathbf{x}_{(t)} + \mathbf{W}_{h,o}\mathbf{h}_{(t-1)} + \mathbf{b}_o)
\end{aligned}
$$

```python
In [2]: print("Done")
```

Done