

# Multinomial Classification: from binary to many classes

What if our targets come from a class  $C$  with more than two discrete values?

The case of  $||C|| > 2$  is called **Multinomial** or **Multiclass** Classification.

Some models (e.g. Decision Trees) can handle Multinomial classification directly.

For those that can't, there are two general approaches to multinomial classification

- Turn the classification task into multiple *binary* classification tasks
  - One versus All others, One versus One
- Generalize the loss function to directly accommodate multiple classes

Both approaches can be viewed as representing target  $\mathbf{y}^{(i)}$  via One Hot Encoding

Notice that the target  $\mathbf{y}$  and predictions  $\hat{\mathbf{y}}$  are now **vectors** of length greater than 1.

## Prediction as probability

Because the representations of  $\mathbf{y}^{(i)}$  and  $\hat{\mathbf{y}}^{(i)}$

- are of length  $\|C\|$
- have elements in the range  $[0, 1]$
- and whose elements sum to 1

both  $\mathbf{y}^{(i)}$  and  $\hat{\mathbf{y}}^{(i)}$  can be viewed as *probability distributions* over  $\|C\|$  discrete values.

- Target  $\mathbf{y}^{(i)}$  has all the probability lumped at a single value
- Prediction  $\hat{\mathbf{y}}^{(i)}$  may *spread* the probability across multiple values
  - $\hat{\mathbf{y}}_j^{(i)}$  is the probability that the target is  $c_j$
  - There may be multiple class values with non-zero probability
  - Usually choose the  $c_k$  with largest probability as the single prediction, if required

$$\operatorname{argmax}_k \hat{\mathbf{y}}_k$$

# Multinomial classification using multiple binary classifiers

## One versus all

The One versus All (OvA) method creates  $||C||$  binary classifiers

- One for each  $c \in C$
- The classifier for class  $c$  identifies
  - Positive examples as those having target  $c$
  - Negative examples as those having targets other than  $c$

For the binary classifier for class  $c$ , let

- $\hat{p}^c(\mathbf{x})$  denote the prediction of example  $\mathbf{x}$  being Positive (i.e., class  $c$ ) made by this binary classifier

Combining the predictions for each class into a vector  $\hat{p}$  of length  $||C||$  such that

$$\hat{p}_c = \hat{p}^c(\mathbf{x})$$

Note that the elements of  $\hat{p}$  may not sum to 1, so in order to create a probability vector we need to normalize its elements in order to create the OvA prediction vector

$$\hat{y}_c(\mathbf{x}) = \frac{\hat{p}^c(\mathbf{x})}{\sum_{c' \in C} \hat{p}^{c'}(\mathbf{x})}$$

That is: it normalizes the probabilities so that they sum to 1 for each example.

### **Note**

We have abused notation by using class  $c$  as a subscript of  $\hat{y}$ ,  $\hat{p}$  rather than the integer  $j$ , where  $c$  is the  $j^{th}$  class in  $C$ .



Note that the binary classifier for each class  $c$  has its own parameters  $\Theta_c$ .

So the number of parameters in the  $\Theta$  for the OvA classifier is  $\|C\|$  times as big as the number of parameters for a single classifier.

Let's be clear on the number of coefficients estimated in One versus All:

For the digit classification problem where there are  $C = 10$  classes the number of parameters is *10 times* that of a binary classifier.

Fortunately, `sklearn` hides all of this from you.

What you *should* realize is that  $||C||$  models are being fit, each with its own parameters.

## One versus one

The One versus One (OvO) method creates  $\frac{||C||*(||C||-1)}{2}$  binary classifiers

- one for each pair  $c, c'$  of distinct values in  $C$
- the classifier for pair  $c, c'$  identifies
  - Positive examples as those having target  $c$
  - Negative examples as those having targets  $c'$

Essentially, OvO creates a "competition" between pairs of classes for a given example  $\mathbf{x}$

- the class that "wins" most often is chosen as the predicted class for the OvO classifier on example  $\mathbf{x}$

# Softmax

A number of binary classifiers (e.g., Logistic Regression)

- Produce a score
- Which is then converted into a probability

For the One Versus All multinomial classification method

- We convert the score (for class  $c_k$ ) into a probability
- We then normalize (across each of the  $||C||$  values for  $k$ ) the probabilities

We can go directly from the score (for class  $c_k$ ) to a normalized probability using the Softmax function

- Multinomial generalization of the Sigmoid function

For the binary classifier for class  $c$ , let

- $\hat{s}^c(\mathbf{x})$  denote the score of example  $\mathbf{x}$  produced by this binary classifier

The probability vector  $\hat{\mathbf{y}}$  can be computed by the *Softmax* function

$$\hat{y}_c(\mathbf{x}) = \frac{\exp(s^c(\mathbf{x}))}{\sum_{c \in C} \exp(s^c(\mathbf{x}))}$$

You can see that each  $\hat{y}_c(\mathbf{x}) \in [0, 1]$  and that  $\sum_{c \in C} \hat{y}_c(\mathbf{x}) = 1$  so  $\hat{\mathbf{y}}$  is indeed a probability vector.

By exponentiating the score, the softmax magnifies small differences in scores into larger difference in probability.

To illustrate: suppose we have two relatively close scores  $s^c, s^{c'}$  such that

$$\frac{s^c}{s^{c'}} = M \approx 1$$

- If we normalize scores by dividing a score by the sum (across all scores)
  - $\frac{\hat{y}_c}{\hat{y}_{c'}} = M$
- If we normalize by softmax
  - $\frac{\hat{y}_c}{\hat{y}_{c'}} = \frac{\exp(M\hat{s}_c)}{\exp(\hat{s}^c)} = \exp(\hat{s}_c(M - 1))$

Softmax is most often seen in the context of Logistic Regression.



# Multinomial classification by generalizing the loss function

We will deal with the loss functions, both for Binary and Multinomial Classification in a separate module.

- For Binary Classification: the loss function is called Binary Cross Entropy
- The generalization of the loss function to Multinomial Classification is called *Cross Entropy*

## Prediction for multinomial classification

Both approaches create a prediction vector  $\hat{\mathbf{y}}$  that is a probability distribution.

If we need to choose a single target as our prediction, we can choose the one with greatest probability. We can choose the class  $c$  with the largest value in  $\hat{\mathbf{p}}$  as our prediction

$$\operatorname{argmax}_{c \in \{1, \dots, ||C||\}} \hat{\mathbf{y}}_c$$

# Multinomial classification example: MNIST digit classifier

Remember the digit classifier using KNN from our introductory lecture ?

We criticized the model as being one of excessive template matching: one template per training example.

We can now use Logistic Regression to obtain a classifier with *many* fewer parameters.

It will also have the benefit of helping us *interpret how* the classifier is making its predictions.


We won't go into interpretation until a later lecture, but for now: a preview of coming attractions.

Let's fetch the data and visualize it.

```
In [5]: mnh.setup()  
        mnh.visualize()
```

Retrieving MNIST\_784 from cache

Out[5]:





```
In [6]: print("Training set: X shape={xs}, y shape: {ys}".format(xs=mnh.X_train.shape, y
s=mnh.y_train.shape) )
print("Training labels: y is of type {t}".format(t=type(mnh.y_train[0])) ) )
```

```
Training set: X shape=(5000, 784), y shape: (5000,)
Training labels: y is of type <class 'str'>
```

The training set  $\mathbf{X}$  consists of 5000 examples, each having 784 features.

The 784 features are pixel intensity values (1=white, 0=black), visualized as a  $(28 \times 28)$  image.

Importantly, the labels (targets) are strings, i.e, string "0" rather than integer 0.

$$C = \{ "0", "1", \dots, "9" \}$$

Let's fit a Logistic Regression model.



```
In [7]: mnist_lr = mnh.fit()
```

How did we do, i.e., what was the Performance Metric?

```
In [8]: clf = mnh.clf
score = clf.score(mnh.X_test, mnh.y_test)

# How many zero coefficients were forced by the penalty ?
sparsity = np.mean(clf.coef_ == 0) * 100

print("Test score with {p} penalty:{s:.2f}".format(p=clf.penalty, s=score) )
print("Sparsity with {p} penalty: {s:.2f}.".format(p=clf.penalty, s=sparsity) )
```

```
Test score with l2 penalty:0.87
Sparsity with l2 penalty: 16.07.
```

We achieved an accuracy on the Test set of about 88%.

Is this good ? We'll probe that question in a later lecture.

For now: it sounds pretty good, but

- In a Test set with equal quantities of each digit
- We could get *all* instances of a single digit wrong and still achieve 90% accuracy !
- **Lesson:** absolute numbers are misleading

Also notice that `LogisticRegression` used an L2 penalty (Ridge Regression)

- That caused about 16% of the parameters to become 0.

How many parameters did we fit (i.e., what is the size of  $\Theta$ ) ?

```
In [9]: print("The classifier non-intercept parameters shape: {nc}; intercept parameter  
s shape: {ni}".format(  
    nc=mnh.clf.coef_.shape,  
    ni=mnh.clf.intercept_.shape  
)  
    )
```

The classifier non-intercept parameters shape: (10, 784); intercept parameter  
s shape: (10,)

sklearn separately stores

- the intercept (`clf.intercept_`): the parameter associated with the constant column in  $\mathbf{X}'$ )
- all other parameters (`clf.coef_`)



As you can see from the leading dimension (10) there are essentially  $\|C\|$  binary classifiers

- One parameter per element of the feature vector
- Plus one intercept/constant parameter

In total  $\Theta$  has  $10 * (784 + 1) = 7850$  parameters.

More precisely

- The target vector  $\mathbf{y}$  is of length  $\|C\| = 10$ , i.e., OHE target
  - We have previously only seen scalar targets
- `LogisticRegression` is performing One versus All (OvA) classification
- Because  $\|\mathbf{y}^{(i)}\| > 1$ , it is using a Cross Entropy Loss in the Loss function

Compare this to the KNN classifier from the first lecture

- one template per example, at  $(28 \times 28) = 784$  parameters per example
- times  $m = 5000$  examples

So the Logistic Classifier uses about  $m = 5000$  times fewer parameters.

What do the 784 non-intercept parameters look like ?

That is: what is the "template" for each class (digit) ?


Since there is one parameter per pixel, ordered in the same way as the input image pixels

- We can display the 784 parameters as a  $(28 \times 28)$  image.

Remember: there is one parameter vector (template) for each of the  $||C|| = 10$  classes.

```
In [10]: mnist_fig, mnist_ax = mnh.plot_coeff()
```

Parameters for...



Our model learned a template, per digit, which hopefully captures the "essence" of the digit

- Fuzzy, since it needs to match many possible examples of the digit, each written differently

We will "interpret" these coefficients in a subsequent lecture but, for now:

- Dark colored parameters indicate the template for the pixel best matches dark input pixels
- Bright colored parameters indicate the template for the pixel best matches bright input pixels

So the "essence" of an image representing the "1" digit is a vertical band of bright pixels.

**TIP** The `fetch_mnist_784` routine in the module takes a **long** time to execute. Caching results makes you more productive.

```
In [11]: print("Done")
```

Done